

---

# **Remote Sensing Image Analysis with R**

**Aniruddah Ghosh and Robert J. Hijmans**

**May 03, 2019**



# CONTENTS

<b>1</b>	<b>Introduction</b>	<b>3</b>
1.1	Terminology . . . . .	3
1.2	Data . . . . .	4
1.3	Resources . . . . .	4
1.4	R packages . . . . .	4
<b>2</b>	<b>Exploration</b>	<b>5</b>
2.1	View image properties . . . . .	5
2.2	Image information and statistics . . . . .	6
2.3	Visualize single and multi-band imagery . . . . .	7
2.4	Subset and rename spectral bands . . . . .	10
2.5	Spatial subset or crop . . . . .	11
2.6	Saving results to disk . . . . .	11
2.7	Relation between bands . . . . .	11
2.8	Extract raster values . . . . .	13
2.9	Spectral profiles . . . . .	14
<b>3</b>	<b>Basic mathematical operations</b>	<b>17</b>
3.1	Vegetation indices . . . . .	17
3.2	Histogram . . . . .	18
3.3	Thresholding . . . . .	19
3.4	Principal component analysis . . . . .	23
<b>4</b>	<b>Unsupervised Image Classification</b>	<b>29</b>
4.1	kmeans classification . . . . .	29
<b>5</b>	<b>Supervised Image Classification</b>	<b>33</b>
5.1	Reference data . . . . .	33
5.2	Generate sample sites . . . . .	34
5.3	Extract cell values for the sample sites . . . . .	35
5.4	Train the classifier using training samples . . . . .	36
5.5	Classify . . . . .	37
5.6	Model evaluation . . . . .	39



Aniruddha Ghosh and Robert J. Hijmans



## INTRODUCTION

This book provides a short introduction to satellite data analysis with R. Before reading this you should first learn the [basics of the raster package](#).

Getting satellite images for a specific project remains a challenging task. You have to find data that is suitable for your objectives, and that you can get access to. Important properties to consider while searching the remotely sensed (satellite) data include:

1. [Spatial resolution](#)
2. [Temporal resolution](#), including the return time; availability of historical images, and for a particular moment in time)
3. [Spectral resolution](#), that is, the parts of the electromagnetic spectrum (wavelengths) that are sensed
4. Radiometric resolution (sensor sensitivity; ability to measure small differences)
5. Quality (e.g. the presence of cloud-cover or of artefacts in the data (read about problems in [Landsat ETM+](#)))

There are numerous sources of remotely sensed data from satellites. Generally, the very high spatial resolution data is available as (costly) commercial products. Lower spatial resolution data is freely available from NASA, ESA and other missions. In this tutorial we'll use freely available [Landsat 8](#), [Landsat 7](#), [Landsat 5](#), [Sentinel](#) and [MODIS](#) data. The [Landsat program](#) started in 1972 and is the longest running Earth-observation satellite program.

You can access public satellite data from several sources, including:

- i. <http://earthexplorer.usgs.gov/>
- ii. <https://lpdaacsvc.cr.usgs.gov/appeears/>
- iii. <https://search.earthdata.nasa.gov/search>
- iv. [https://lpdaac.usgs.gov/data\\_access/data\\_pool](https://lpdaac.usgs.gov/data_access/data_pool)
- v. <https://scihub.copernicus.eu/>
- vi. <https://aws.amazon.com/public-data-sets/landsat/>

This [web site](#) lists 15 sources of freely available remote sensing data.

It is possible to download some satellite data using R-packages. You can use the [MODIS](#) or [MODISTools](#) package to search, download and pre-process different [MODIS products](#).

### 1.1 Terminology

Most remote sensing products consist of observations of reflectance data. That is, they are measures of the intensity of the sun's radiation that is reflected by the earth. Reflectance is normally measured for different wavelengths of the electromagnetic spectrum. For example, it can be measured in the red, green, and blue wavelengths. If that is the case, a satellite image can be referred to as a "multi-spectral" (or hyper-spectral if there are many separate wavelengths).

Therefore, a single ‘satellite image’ typically consists of a number of separate raster layers. In remote sensing jargon, these layers are referred to as “bands” (shorthand for “bandwidth”). Other commonly used jargon in remote sensing circles is “pixel”, which is equivalent to “grid cell”.

## 1.2 Data

You can download all the data required for the examples used in this book [here](#). Unzip the file contents and save the data to the R working directory of your computer.

You can also use the below script to download the data.

```
dir.create('data', showWarnings = FALSE)
if (!file.exists('data/rs/samples.rds')) {
  download.file('https://biogeo.ucdavis.edu/data/rsdata/rsdata.zip', dest = 'data/
  ↳rsdata.zip')
  unzip('data/rsdata.zip', exdir='data')
}
```

## 1.3 Resources

Here is a short list of some good resources to learn about remote sensing image analysis:

- Remote Sensing Digital Image Analysis
- Introductory Digital Image Processing: A Remote Sensing Perspective
- A survey of image classification methods and techniques for improving classification performance
- A Review of Modern Approaches to Classification of Remote Sensing Data
- Online remote sensing course

## 1.4 R packages

Here is a list of some R packages for analyzing remote sensing data:

- RStoolbox
- landsat
- hsdar
- rasterVis for visualization



## EXPLORATION

In this chapter we describe how to access and explore attributes of remote sensing images with R. We also show how to plot (make maps).

We will primarily use a spatial subset of Landsat 8 scene collected on June 14, 2017. The subset covers the area between Concord and Stockton, in California, USA.

All Landsat image scenes have a unique product ID and metadata. You can find the information on Landsat sensor, satellite, [location on Earth \(WRS path, WRS row\)](#) and acquisition date from the product ID. For example, the product identifier of the data we will use is 'LC08\_044034\_20170614'. Based on [this guide](#), you can see that the Sensor-Satellite is OLI/TIRS combined Landsat 8, WRS Path 44, WRS Row 34 and collected on June 14, 2017. Landsat scenes are most commonly delivered as zipped file, which contains separate files for the reflectance values at each bandwidth (wavelength).

We will start by exploring and visualizing the data (See [chapter 1] for data downloading instructions if you have not already done so).

### 2.1 View image properties

Create RasterLayer objects for single Landsat layers (bands)

```
library(raster)
## Loading required package: sp

# Blue
b2 <- raster('data/rs/LC08_044034_20170614_B2.tif')

# Green
b3 <- raster('data/rs/LC08_044034_20170614_B3.tif')

# Red
b4 <- raster('data/rs/LC08_044034_20170614_B4.tif')

# Near Infrared (NIR)
b5 <- raster('data/rs/LC08_044034_20170614_B5.tif')
```

Print the variables to check. E.g.

```
b2
## class      : RasterLayer
## dimensions : 1245, 1497, 1863765  (nrow, ncol, ncell)
## resolution : 30, 30  (x, y)
## extent     : 594090, 639000, 4190190, 4227540  (xmin, xmax, ymin, ymax)
## crs        : +proj=utm +zone=10 +datum=WGS84 +units=m +no_defs +ellps=WGS84_
               ↪ +towgs84=0,0,0
```

(continues on next page)

(continued from previous page)

```
## source      : c:/github/rspatial/rspatial-web/source/rs/_R/data/rs/LC08_044034_
↪20170614_B2.tif
## names       : LC08_044034_20170614_B2
## values      : 0.0748399, 0.7177562 (min, max)
```

You can see the spatial resolution, extent, number of layers, coordinate reference system and more.

## 2.2 Image information and statistics

The below shows how you can access various properties from a Raster\* object (this is the same for any raster data set).

```
# coordinate reference system (CRS)
crs(b2)
## CRS arguments:
## +proj=utm +zone=10 +datum=WGS84 +units=m +no_defs +ellps=WGS84
## +towgs84=0,0,0

# Number of rows, columns, or cells
ncell(b2)
## [1] 1863765
dim(b2)
## [1] 1245 1497 1

# spatial resolution
res(b2)
## [1] 30 30

# Number of bands
nlayers(b2)
## [1] 1

# Do the bands have the same extent, number of rows and columns, projection,
↪resolution, and origin
compareRaster(b2,b3)
## [1] TRUE
```

You can create a RasterStack (an object with multiple layers) from the existing RasterLayer (single band) objects.

```
landsatRGB <- stack(b4, b3, b2)
landsatFCC <- stack(b5, b4, b3)

# Check the properties of the RasterStack
landsatRGB
## class      : RasterStack
## dimensions : 1245, 1497, 1863765, 3 (nrow, ncol, ncell, nlayers)
## resolution : 30, 30 (x, y)
## extent      : 594090, 639000, 4190190, 4227540 (xmin, xmax, ymin, ymax)
## crs         : +proj=utm +zone=10 +datum=WGS84 +units=m +no_defs +ellps=WGS84
↪+towgs84=0,0,0
## names       : LC08_044034_20170614_B4, LC08_044034_20170614_B3, LC08_044034_
↪20170614_B2
## min values  : 0.02084067, 0.04259216, 0.
↪07483990
## max values  : 0.7861769, 0.6924697, 0.
↪7177562
```

Notice the order of the layers. These are suitable for plotting different color composites. You can learn more about color composites in remote sensing [here](#) and also in the section below.

You can also create the RasterStack using the filenames.

```
# first create a list of raster layers to use
filenames <- paste0('data/rs/LC08_044034_20170614_B', 1:11, ".tif")
filenames
## [1] "data/rs/LC08_044034_20170614_B1.tif"
## [2] "data/rs/LC08_044034_20170614_B2.tif"
## [3] "data/rs/LC08_044034_20170614_B3.tif"
## [4] "data/rs/LC08_044034_20170614_B4.tif"
## [5] "data/rs/LC08_044034_20170614_B5.tif"
## [6] "data/rs/LC08_044034_20170614_B6.tif"
## [7] "data/rs/LC08_044034_20170614_B7.tif"
## [8] "data/rs/LC08_044034_20170614_B8.tif"
## [9] "data/rs/LC08_044034_20170614_B9.tif"
## [10] "data/rs/LC08_044034_20170614_B10.tif"
## [11] "data/rs/LC08_044034_20170614_B11.tif"

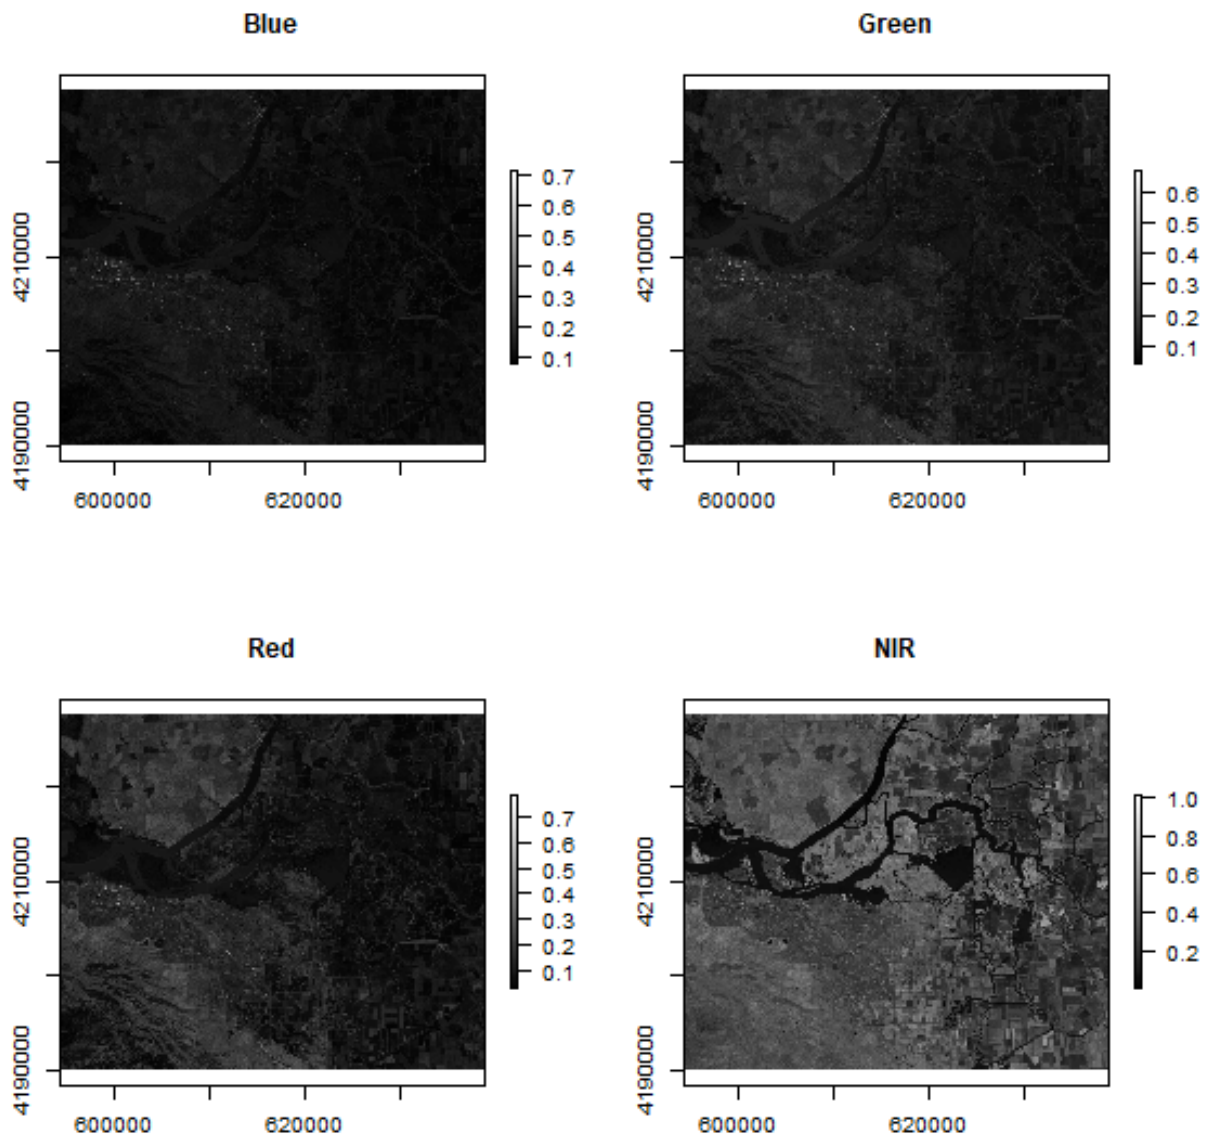
landsat <- stack(filenames)
landsat
## class      : RasterStack
## dimensions : 1245, 1497, 1863765, 11  (nrow, ncol, ncell, nlayers)
## resolution : 30, 30  (x, y)
## extent      : 594090, 639000, 4190190, 4227540  (xmin, xmax, ymin, ymax)
## crs         : +proj=utm +zone=10 +datum=WGS84 +units=m +no_defs +ellps=WGS84_
↪ +towgs84=0,0,0
## names       : LC08_044034_20170614_B1, LC08_044034_20170614_B2, LC08_044034_
↪ 20170614_B3, LC08_044034_20170614_B4, LC08_044034_20170614_B5, LC08_044034_20170614_
↪ B6, LC08_044034_20170614_B7, LC08_044034_20170614_B8, LC08_044034_20170614_B9, LC08_
↪ 044034_20170614_B10, LC08_044034_20170614_B11
## min values :          9.641791e-02,          7.483990e-02,          4.
↪ 259216e-02,          2.084067e-02,          8.457669e-04,          -7.872183e-
↪ 03,          -5.052945e-03,          3.931751e-02,          -4.337332e-04,
↪          2.897978e+02,          2.885000e+02
## max values :          0.73462820,          0.71775615,          0.
↪ 69246972,          0.78617686,          1.01243150,          1.04320455,
↪          1.11793602,          0.82673049,          0.03547901,
↪          322.43139648,          317.99530029
```

Above we created a RasterStack with 11 layers. The layers represent reflection intensity in the following wavelengths: Ultra Blue, Blue, Green, Red, Near Infrared (NIR), Shortwave Infrared (SWIR) 1, Shortwave Infrared (SWIR) 2, Panchromatic, Cirrus, Thermal Infrared (TIRS) 1, Thermal Infrared (TIRS) 2. We won't use the last four layers and we will learn how to remove those in following sections.

## 2.3 Visualize single and multi-band imagery

You can plot individual layers of a RasterStack of a multi-spectral image.

```
par(mfrow = c(2,2))
plot(b2, main = "Blue", col = gray(0:100 / 100))
plot(b3, main = "Green", col = gray(0:100 / 100))
plot(b4, main = "Red", col = gray(0:100 / 100))
plot(b5, main = "NIR", col = gray(0:100 / 100))
```




Check the legends. They represent the range of values in each layer, and they range from 0 and 1. Notice the difference in shading and range of legends between the different bands. This is because different Earth surface features reflect the incident solar radiation differently. Each layer represent how much incident solar radiation is reflected for that particular wavelength. For example, vegetation reflects more energy in NIR than other wavelengths and thus appears brighter in NIR wavelength. However water absorbs most of the incident energy in NIR regions and appears dark.

However we did not get much information from these grey-scale plots, but these are often combined to create more interesting plots. To combine three bands, we can use `plotRGB`. To make a “true/natural color” image (that is, something that looks like a normal photograph), we need to select the bands that we want to render in the red, green and blue regions. For this Landsat image, `r = 3` (red), `g = 2` (green), `b = 1` (blue) will plot the true color composite (vegetation in green, water blue etc). You can also supply additional arguments to `plotRGB` to improve the visualization (e.g. a linear stretch of the values, using `stretch = "lin"`).

```
plotRGB(landsatRGB, axes = TRUE, stretch = "lin", main = "Landsat True Color Composite")
```

(continues on next page)

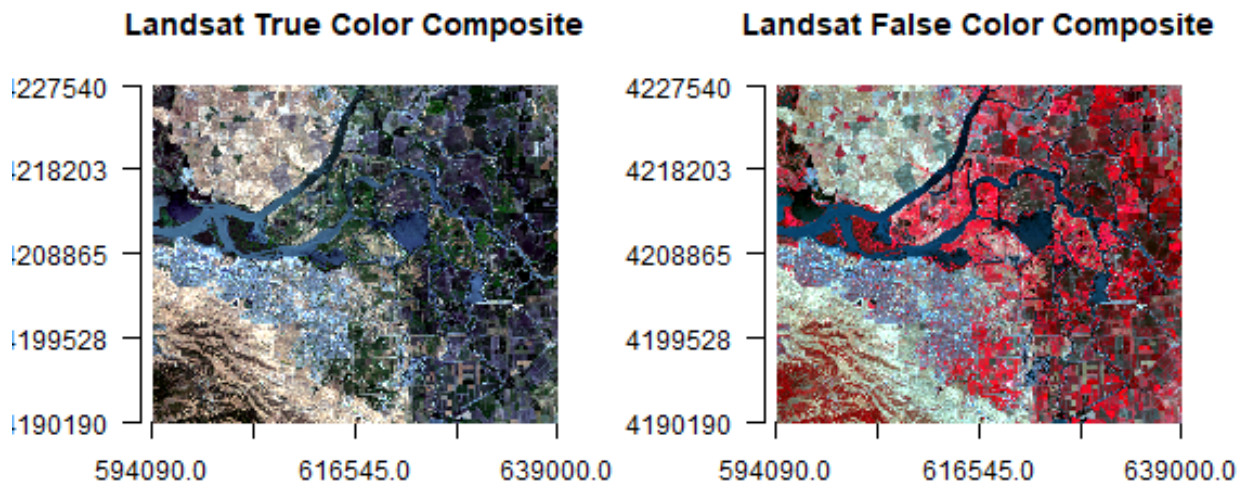
(continued from previous page)



This RGB-plot has lot more information than the previous ones, although it has been created using the same bands. Another popular image visualization method in remote sensing is known “false color” image where red, green and blue bands are replaced by other bands. Selecting  $r = \text{NIR}$ ,  $g = \text{red}$ ,  $b = \text{green}$  will plot a “false color” composite. This representation is popular as it makes it easy to see the vegetation (in red).

```
par(mfrow = c(1,2))
plotRGB(landsatRGB, r=1, g=2, b=3, axes=TRUE, stretch="lin", main="Landsat True Color_
↪Composite")
plotRGB(landsatFCC, r=1, g=2, b=3, axes=TRUE, stretch="lin", main="Landsat False_
↪Color Composite")
```





*Note:* Always check for package documentation (`help(plotRGB)`) for other arguments that can be added (like `scale`) to improve or modify the image.

**Exercise 1** Use the RasterStack `landsat` to create a true and false color composite (*hint* remember the position of the bands in the stack).

## 2.4 Subset and rename spectral bands

You can select specific layers (bands) using `subset` function, or via indexing.

```
# select first 3 bands only
landsatsub1 <- subset(landsat, 1:3)
# same
landsatsub2 <- landsat[[1:3]]

# Number of bands in original and new data
nlayers(landsat)
## [1] 11
nlayers(landsatsub1)
## [1] 3
nlayers(landsatsub2)
## [1] 3
```

As mentioned above, we have no use here for the last four bands in `landsat`. You can remove those using

```
landsat <- subset(landsat, 1:7)
```

Set the names of the bands using the following:

```
names(landsat)
## [1] "LC08_044034_20170614_B1" "LC08_044034_20170614_B2"
## [3] "LC08_044034_20170614_B3" "LC08_044034_20170614_B4"
## [5] "LC08_044034_20170614_B5" "LC08_044034_20170614_B6"
## [7] "LC08_044034_20170614_B7"
names(landsat) <- c('ultra-blue', 'blue', 'green', 'red', 'NIR', 'SWIR1', 'SWIR2')
```

(continues on next page)

(continued from previous page)

```
names(landsat)
## [1] "ultra.blue" "blue"      "green"      "red"        "NIR"
## [6] "SWIR1"      "SWIR2"
```

## 2.5 Spatial subset or crop

Spatial subsetting can be used to limit analysis to a geographic subset of the image. Spatial subsets can be created with the `crop` function, using an extent object, or another spatial object from which an Extent can be extracted/created,.

```
# Using extent
extent(landsat)
## class      : Extent
## xmin       : 594090
## xmax       : 639000
## ymin       : 4190190
## ymax       : 4227540
e <- extent(624387, 635752, 4200047, 4210939)

# crop landsat by the extent
landsatcrop <- crop(landsat, e)
```

**Exercise 2** Interactive selection from the image is also possible. Use `drawExtent` and `drawPoly` to select an area of interest. *Note:* drawing will not work for plots generated within the markdown document. Please run the `plot` and `draw` commands from the RStudio console.

**Exercise 3** Use the RasterStack `landsatcrop` to create a true and false color composite (*hint* remember the position of the bands in the stack).

## 2.6 Saving results to disk

At this stage we may want to save the raster to disk using the function `writeRaster`. Multiple file types are supported. We will use the commonly used GeoTiff format. While the layer order is preserved, layer names are unfortunately lost in the GeoTiff format.

```
writeRaster(landsatcrop, filename = "cropped-landsat.tif", overwrite = TRUE)
```

To keep the layer names, you can use the 'raster-grd' format:

```
writeRaster(landsatcrop, filename = "cropped-landsat.grd", overwrite = TRUE)
```

The disadvantage of this format is that not many other programs can read the data, in contrast to the GeoTiff format.


*Note:* Check for package documentation (`help(writeRaster)`) for additional helpful arguments that can be added.

## 2.7 Relation between bands

A scatterplot matrix can be helpful in exploring relationships between raster layers. This can be done with the `pairs()` function of the raster package.

Plot of reflection in the ultra-blue wavelength against reflection in the blue wavelength.

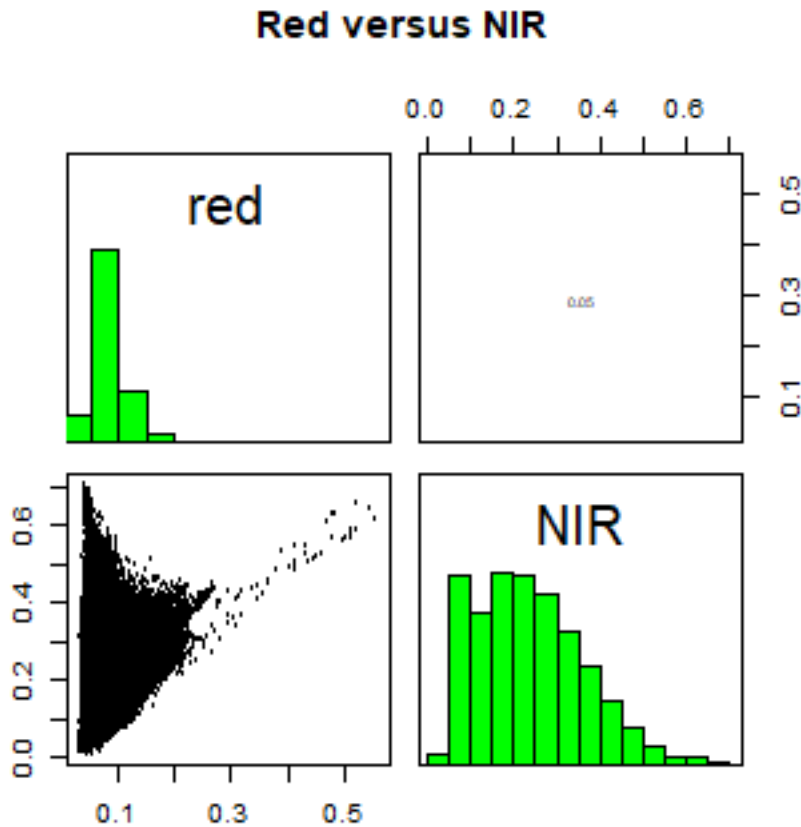
```
pairs(landsatcrop[[1:2]], main = "Ultra-blue versus Blue")
```



Plot of reflection in the red wavelength against reflection in the NIR wavelength.

```
pairs(landsatcrop[[4:5]], main = "Red versus NIR")
```





The first plot reveals high correlations between the blue wavelength regions. Because of the high correlation, we can just use one of the Blue bands without losing much information.

This distribution of points in second plot (between NIR and Red) is unique due to its triangular shape. Vegetation reflects very highly in the NIR range than red and creates the upper corner close to NIR (y) axis. Water absorbs energy from all the bands and occupies the location close to origin. The furthest corner is created due to highly reflecting surface features like bright soil or concrete.

## 2.8 Extract raster values

Often we want to get the values of raster cells for specific geographic locations or area. The `extract` function is used to get raster values at the locations of other spatial data. You can use points, lines, polygons or an Extent (rectangle) object. You can also use cell numbers to extract values. When using points, `extract` returns the values of a `Raster*` object for the cells in which a set of points fall.

```
# load the polygons with land use land cover information
samp <- readRDS('data/rs/samples.rds')

# generate 300 point samples from the polygons;
ptsamp <- spsample(samp, 300, type = 'random')
# add the class information to the point samples from polygons
ptsamp$class <- over(ptsamp, samp)$class
```

(continues on next page)

(continued from previous page)

```
# extract values with points
df <- extract(landsat, ptsamp)

# To see some of the reflectance values
head(df)
##      ultra.blue      blue      green      red      NIR      SWIR1
## [1,]  0.1379691 0.1290126 0.12775482 0.15560021 0.22853172 0.27767318
## [2,]  0.1101888 0.0879602 0.08026151 0.04788366 0.56883532 0.11875495
## [3,]  0.1352366 0.1306825 0.14317387 0.19207680 0.38185486 0.33379769
## [4,]  0.2017490 0.2012068 0.21020667 0.23414850 0.27179617 0.28823447
## [5,]  0.1333065 0.1141357 0.09225412 0.07171705 0.05042097 0.03623804
## [6,]  0.1105358 0.0874831 0.08321087 0.04701620 0.59245187 0.15438578
##
##      SWIR2
## [1,] 0.19363822
## [2,] 0.04150784
## [3,] 0.18745759
## [4,] 0.24750735
## [5,] 0.02932006
## [6,] 0.05449802
```

## 2.9 Spectral profiles

A plot of the spectrum (all bands) for pixels representing a certain earth surface features (e.g. water) is known as a spectral profile. Such profiles demonstrate the differences in spectral properties of various earth surface features and constitute the basis for image analysis. Spectral values can be extracted from any multispectral data set using `extract` function. In the above example, we extracted values of Landsat data for the samples. These samples include: cropland, water, fallow, built and open. First we compute the mean reflectance values for each class and each band.

```
ms <- aggregate(df, list(ptsamp$class), mean)

# instead of the first column, we use rownames
rownames(ms) <- ms[,1]
ms <- ms[,-1]
ms
##      ultra.blue      blue      green      red      NIR      SWIR1
## built      0.1787321 0.17204419 0.17406668 0.19112260 0.24847852 0.24790147
## cropland    0.1124401 0.09043855 0.08602197 0.05467287 0.48582778 0.15340649
## fallow      0.1331698 0.11651040 0.10211767 0.10796453 0.16892134 0.22391816
## open        0.1391130 0.13724872 0.15152051 0.20347356 0.33900667 0.35123016
## water       0.1340993 0.11716628 0.09994984 0.07940147 0.04941895 0.03397931
##
##      SWIR2
## built      0.20483975
## cropland    0.06696569
## fallow      0.19192028
## open        0.21178264
## water       0.02748395
```

Now we will plot the mean spectra of these features.

```
# Create a vector of color for the land cover classes for use in plotting
mycolor <- c('darkred', 'yellow', 'burlywood', 'cyan', 'blue')

#transform ms from a data.frame to a matrix
ms <- as.matrix(ms)
```

(continues on next page)


(continued from previous page)

```
# First create an empty plot
plot(0, ylim=c(0,0.6), xlim = c(1,7), type='n', xlab="Bands", ylab = "Reflectance")

# add the different classes
for (i in 1:nrow(ms)){
  lines(ms[i,], type = "l", lwd = 3, lty = 1, col = mycolor[i])
}

# Title
title(main="Spectral Profile from Landsat", font.main = 2)

# Legend
legend("topleft", rownames(ms),
      cex=0.8, col=mycolor, lty = 1, lwd =3, bty = "n")
```



The spectral profile shows (dis)similarity in the reflectance of different features on the earth's surface (or above it). Unsurprisingly, the spectral signatures of 'crop' and 'vegetation' are similar. 'Water' shows relatively low reflection in all wavelengths, and 'built', 'fallow' and 'open' have relatively high reflection in the longer wavelengths.



## BASIC MATHEMATICAL OPERATIONS

The `raster` package supports many mathematical operations. Math operations are generally performed per pixel. First we will learn about basic arithmetic operations on bands. First example is a custom math function that calculates the Normalized Difference Vegetation Index (NDVI). Learn more about [vegetation indices here](#) and [NDVI](#).

We will use the same Landsat data as in Chapter 2.

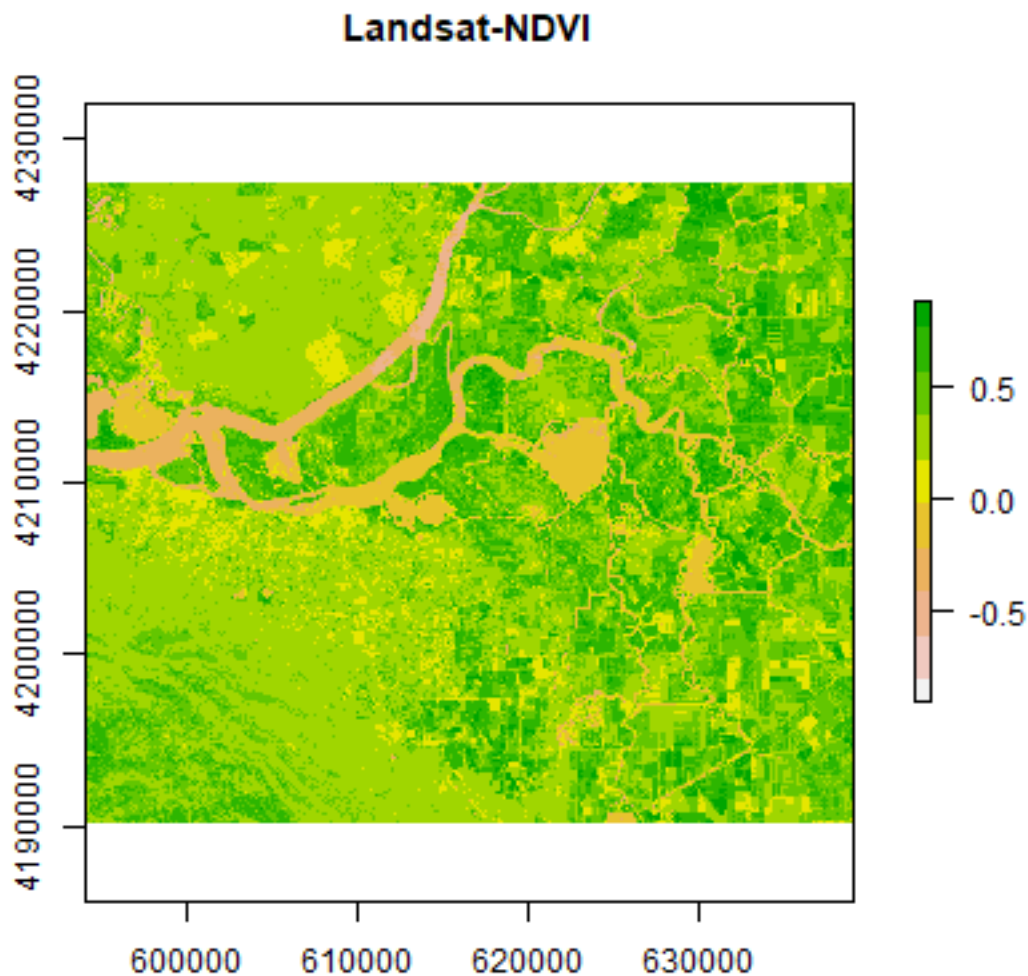
```
library(raster)
raslist <- paste0('data/rs/LC08_044034_20170614_B', 1:11, ".tif")
landsat <- stack(raslist)
landsatRGB <- landsat[[c(4,3,2)]]
landsatFCC <- landsat[[c(5,4,3)]]
```

### 3.1 Vegetation indices

Let's define a general function for a ratio based (vegetation) index. In the function below, `img` is a multilayer `Raster*` object and `i` and `k` are the indices of the layers to compute the vegetation index.

```
VI <- function(img, k, i) {
  bk <- img[[k]]
  bi <- img[[i]]
  vi <- (bk - bi) / (bk + bi)
  return(vi)
}
```

```
# For Landsat NIR = 5, red = 4.
ndvi <- VI(landsat, 5, 4)
plot(ndvi, col = rev(terrain.colors(10)), main = 'Landsat-NDVI')
```



You can see the variation in greenness from the plot.

**Exercise 1** Adapt the function to compute indices which will highlight i) water and ii) built-up. Hint: Use the spectral profile plot to find the bands having maximum and minimum reflectance for these two classes.

## 3.2 Histogram


We can explore the distribution of values contained within our raster using the `hist()` function which produces a histogram. Histograms are often useful in identifying outliers and bad data values in our raster data.

```
# view histogram of data
hist(ndvi,
     main = "Distribution of NDVI values",
     xlab = "NDVI",
     ylab = "Frequency",
     col = "wheat",
     xlim = c(-0.5, 1),
```

(continues on next page)

(continued from previous page)

```
breaks = 30,
xaxt = 'n')
axis(side=1, at = seq(-0.5,1, 0.05), labels = seq(-0.5,1, 0.05))
```



We will refer to this histogram for the following sub-section on thresholding.

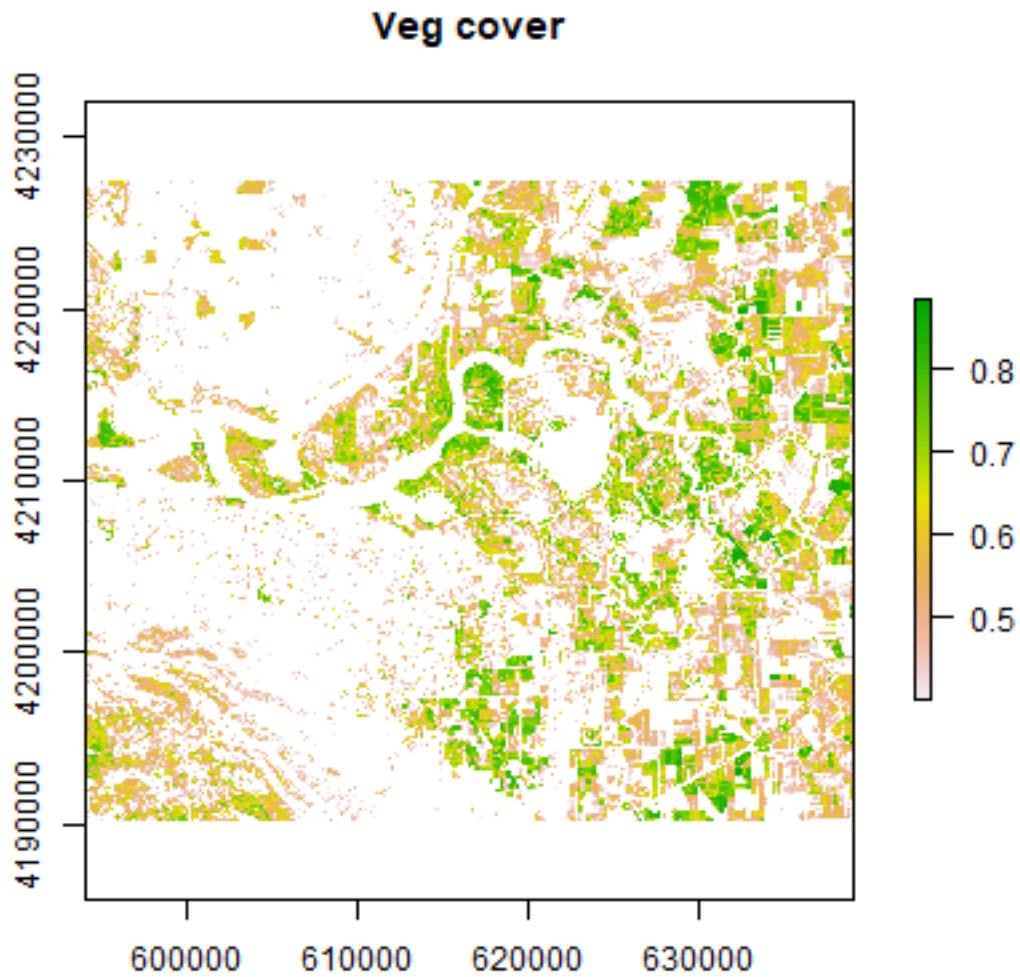
**Exercise 2** Make a histogram of the values of another vegetation index than NDVI.

### 3.3 Thresholding

We can apply basic rules to get an estimate of spatial extent of different Earth surface features. Note that NDVI values are standardized and ranges between -1 to +1. Higher values indicate more green cover.

Cells with NDVI values greater than 0.4 are definitely vegetation. The following operation masks all cells that are perhaps not vegetation.


```
veg <- reclassify(ndvi, cbind(-Inf, 0.4, NA))  
plot(veg, main = 'Veg cover')
```



Let's map the area that corresponds to the peak between 0.25 and 0.3 in the NDVI histogram.


```
land <- reclassify(ndvi, c(-Inf, 0.25, NA, 0.25, 0.3, 1, 0.3, Inf, NA))  
plot(land, main = 'What is it?')
```





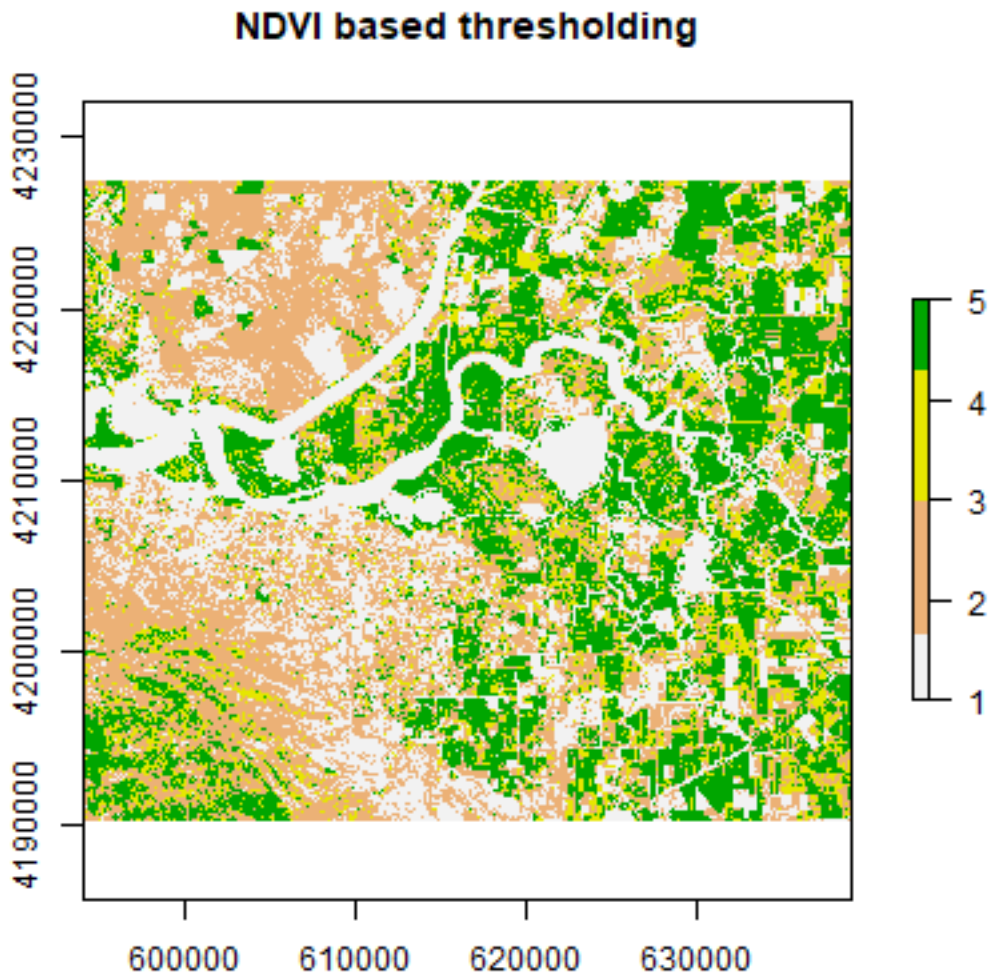
These may be the open areas. You can plot land over original landsatFCC raster to find out more.

```
plotRGB(landsatRGB, r=1, g=2, b=3, axes=TRUE, stretch="lin", main="Landsat False-
  ↳Color Composite")
plot(land, add=TRUE, legend=FALSE)
```



You can also create classes for different amount of vegetation presence.

```
veg1 <- reclassify(ndvi, c(-Inf,0.25,1, 0.25,0.3,2, 0.3,0.4,3, 0.4,0.5,4, 0.5,Inf, 5))  
plot(veg1,col = rev(terrain.colors(4)), main = 'NDVI based thresholding')
```




**Exercise 3** Is it possible to find water using thresholding of NDVI or any other indices?

### 3.4 Principal component analysis

Multi-spectral data are sometimes transformed to help to reduce the dimensionality and noise in the data. The principal components transform is a generic data reduction method that can be used to create a few uncorrelated bands from a larger set of correlated bands.

You can calculate the same number of principal components as the number of input bands. The first principal component (PC) explains the largest percentage of variance and other PCs explain additional the variance in decreasing order.

```
set.seed(1)
sr <- sampleRandom(landsat, 10000)
plot(sr[,c(4,5)], main = "NIR-Red plot")
```



This is known as vegetation and soil-line plot (Same as the scatter plot in earlier section).


**Exercise 4** Can you guess the directions of ‘principal components’ from this scatter plot?

```
pca <- prcomp(sr, scale = TRUE)
pca
## Standard deviations (1, .., p=11):
## [1] 2.52687022 1.40533428 1.09901821 0.92507423 0.53731672 0.42657919
## [7] 0.28002273 0.12466139 0.09031384 0.04761419 0.03609857
##
## Rotation (n x k) = (11 x 11):
##
##          PC1          PC2          PC3          PC4
## LC08_044034_20170614_B1 0.2973469 -0.3516438 -0.29454767 -0.06983456
## LC08_044034_20170614_B2 0.3387629 -0.3301194 -0.16407835 -0.03803678
## LC08_044034_20170614_B3 0.3621173 -0.2641152 0.07373240 0.04090884
## LC08_044034_20170614_B4 0.3684797 -0.1659582 0.10260552 0.03680852
## LC08_044034_20170614_B5 0.1546322 0.1816252 0.71354112 0.32017620
## LC08_044034_20170614_B6 0.3480230 0.2108982 0.23064060 0.16598851
```

(continues on next page)


(continued from previous page)

```
## LC08_044034_20170614_B7 0.3496281 0.2384417 -0.11662258 0.07600209
## LC08_044034_20170614_B8 0.3490464 -0.2007305 0.08765521 0.02303421
## LC08_044034_20170614_B9 0.1314827 0.1047365 0.33741447 -0.92325315
## LC08_044034_20170614_B10 0.2497611 0.4912132 -0.29286315 -0.02950655
## LC08_044034_20170614_B11 0.2472765 0.4931489 -0.29515754 -0.03176549
##
## PC5 PC6 PC7 PC8
## LC08_044034_20170614_B1 0.49474685 0.175510232 -0.23948553 0.215745065
## LC08_044034_20170614_B2 0.22121122 0.094184121 0.06447037 0.216537517
## LC08_044034_20170614_B3 0.08482031 0.009040232 0.30511210 -0.518233675
## LC08_044034_20170614_B4 -0.33835490 -0.066844529 0.60174786 0.012437959
## LC08_044034_20170614_B5 0.51960822 -0.059561476 -0.07348455 -0.083217504
## LC08_044034_20170614_B6 -0.29437062 0.317984598 -0.02106132 0.632178645
## LC08_044034_20170614_B7 -0.25404931 0.525411720 -0.40543545 -0.478543437
## LC08_044034_20170614_B8 -0.31407992 -0.673584139 -0.52642131 0.003527306
## LC08_044034_20170614_B9 0.03040161 0.059642905 -0.03152221 -0.002775800
## LC08_044034_20170614_B10 0.16317572 -0.243735973 0.14341520 0.041736319
## LC08_044034_20170614_B11 0.19294569 -0.241611777 0.11997475 -0.022446494
##
## PC9 PC10 PC11
## LC08_044034_20170614_B1 0.122812108 0.535959306 0.1203473847
## LC08_044034_20170614_B2 0.091063964 -0.773112627 -0.1817872036
## LC08_044034_20170614_B3 -0.644305383 0.070860458 0.0540175730
## LC08_044034_20170614_B4 0.543822097 0.218324141 0.0135097158
## LC08_044034_20170614_B5 0.209682702 -0.040186292 -0.0004965182
## LC08_044034_20170614_B6 -0.397210135 0.089423617 -0.0045305608
## LC08_044034_20170614_B7 0.248871690 -0.074907393 0.0003958921
## LC08_044034_20170614_B8 -0.012642132 0.002411975 0.0007749022
## LC08_044034_20170614_B9 0.003176077 -0.000832654 0.0019986985
## LC08_044034_20170614_B10 -0.003574004 -0.158727672 0.6900281980
## LC08_044034_20170614_B11 -0.043475408 0.148102091 -0.6878990264
screepplot(pca)
```



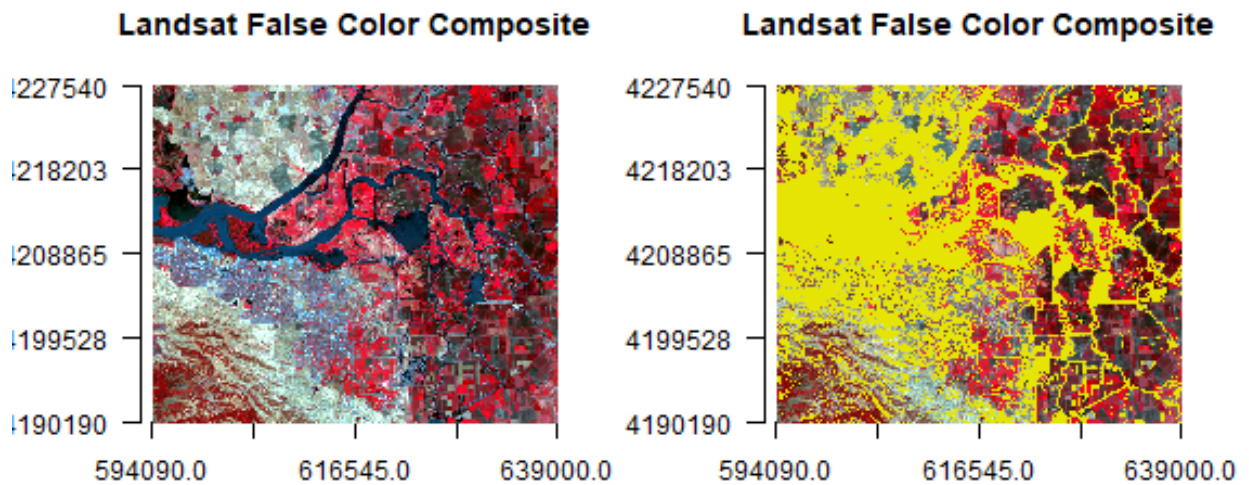
```
pci <- predict(landsat, pca, index = 1:2)
plot(pci[[1]])
```





The first principal component highlights the boundaries between land use classes or spatial details, which is the most common information among all wavelengths. It is difficult to understand what the second principal component is highlighting. Let's try thresholding again:

```
pc2 <- reclassify(pci[[2]], c(-Inf, 0, 1, 0, Inf, NA))
par(mfrow = c(1, 2))
plotRGB(landsatFCC, r = 1, g = 2, b = 3, axes = TRUE, stretch = "lin", main =
  ↪ "Landsat False Color Composite")
plotRGB(landsatFCC, r = 1, g = 2, b = 3, axes = TRUE, stretch = "lin", main =
  ↪ "Landsat False Color Composite")
plot(pc2, legend = FALSE, add = TRUE)
```



**Exercise 5** Can you guess the surface features dominating second principal component? (*hint*: Inspect the unmasked areas; there could be multiple features).

To learn more about the information contained in the vegetation and soil line plots read this paper by [Gitelson et al.](#). An extension of PCA in remote sensing is known as [Tasseled-cap Transformation](#).



## UNSUPERVISED IMAGE CLASSIFICATION

In this chapter we explore unsupervised classification. Various unsupervised classification algorithms exist, and the choice of algorithm can affect the results. We will explore only one algorithm (k-means) to illustrate the general principle.

For this example, we will follow the [National Land Cover Database 2011 \(NLCD 2011\)](#) classification scheme for a subset of the Central Valley regions. You will use cloud-free composite image from [Landsat 5](#) with 6 bands.

```
library(raster)
landsat5 <- stack('data/rs/centralvalley-2011LT5.tif')
names(landsat5) <- c('blue', 'green', 'red', 'NIR', 'SWIR1', 'SWIR2')
```

**Exercise 1** Make a 3-band False Color Composite plot of landsat5.

In unsupervised classification, we don't supply any training data. This particularly is useful when we don't have prior knowledge of the study area. The algorithm groups pixels with similar spectral characteristics into unique clusters/classes/groups following some statistically determined conditions (e.g. minimizing mean square root error within each cluster). You have to re-label and combine these spectral clusters into information classes (for e.g. land-use land-cover). Unsupervised algorithms are often referred as clustering.

To get satisfactory results from unsupervised classification, a good practice is to start with large number of centers (more clusters) and merge/group/recode similar clusters by inspecting the original imagery.

Learn more about K-means and other unsupervised-supervised algorithms [here](#).

We will perform unsupervised classification on a spatial subset of the ndvi layer.

```
ndvi <- (landsat5[['NIR']]-landsat5[['red']])/(landsat5[['NIR']]+landsat5[['red']])
```

We will do kmeans clustering of the ndvi data. First we use `crop` to make a spatial subset of the ndvi, to allow for faster processing (you can select any extent using the `drawExtent()` function).

### 4.1 kmeans classification

```
# Extent to crop ndvi layer
e <- extent(-121.807, -121.725, 38.004, 38.072)

# crop landsat by the extent
ndvi <- crop(ndvi, e)
ndvi
## class      : RasterLayer
## dimensions : 252, 304, 76608 (nrow, ncol, ncell)
## resolution : 0.0002694946, 0.0002694946 (x, y)
## extent     : -121.807, -121.725, 38.00413, 38.07204 (xmin, xmax, ymin, ymax)
```

(continues on next page)

(continued from previous page)

```
## crs      : +proj=longlat +datum=WGS84 +no_defs +ellps=WGS84 +towgs84=0,0,0
## source   : memory
## names     : layer
## values    : -0.3360085, 0.7756007 (min, max)

# convert the raster to vector/matrix
nr <- getValues(ndvi)
str(nr)
## num [1:76608] 0.245 0.236 0.272 0.277 0.277 ...
```

Please note that `getValues` converted the `ndvi` `RasterLayer` to array. Now we will perform the `kmeans` clustering on the array object and inspect the output.

```
# It is important to set the seed generator because `kmeans` initiates the centers in
  ↳ random locations
set.seed(99)

# We want to create 10 clusters, allow 500 iterations, start with 5 random sets using
  ↳ "Lloyd" method
kmncluster <- kmeans(na.omit(nr), centers = 10, iter.max = 500, nstart = 5, algorithm=
  ↳ "Lloyd")

# kmeans returns an object of class "kmeans"
str(kmncluster)
## List of 9
## $ cluster      : int [1:76608] 6 6 3 3 3 3 3 6 6 6 ...
## $ centers       : num [1:10, 1] 0.38986 0.00498 0.29997 -0.14263 -0.20902 ...
## ..- attr(*, "dimnames")=List of 2
## .. ..$ : chr [1:10] "1" "2" "3" "4" ...
## .. ..$ : NULL
## $ totss        : num 6459
## $ withinss     : num [1:10] 5.18 6.13 4.91 6.11 5.75 ...
## $ tot.withinss : num 55.8
## $ betweenss    : num 6403
## $ size         : int [1:10] 8736 4550 7156 8624 11672 6807 8932 5198 5040 9893
## $ iter         : int 221
## $ ifault       : NULL
## - attr(*, "class")= chr "kmeans"
```

`kmeans` returns an object with 9 elements. The length of the `cluster` element within `kmncluster` is 76608 which same as length of `nr` created from the `ndvi`. The cell values of `kmncluster$cluster` range between 1 to 10 corresponding to the input number of cluster we provided in the `kmeans` function. `kmncluster$cluster` indicates the cluster label for corresponding pixel. We need to convert the `kmncluster$cluster` array back to `RasterLayer` of the same dimension as the `ndvi`.

```
# First create a copy of the ndvi layer
knr <- ndvi

# Now replace raster cell values with kmncluster$cluster array
knr[] <- kmncluster$cluster

# Alternative
values(knr) <- kmncluster$cluster
knr
## class      : RasterLayer
## dimensions : 252, 304, 76608 (nrow, ncol, ncell)
```

(continues on next page)

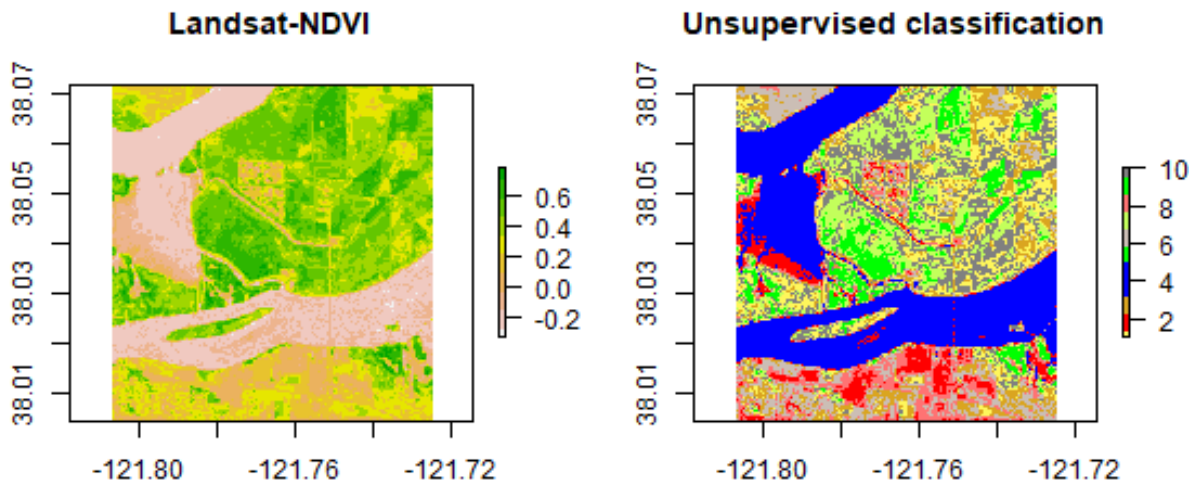
(continued from previous page)

```
## resolution : 0.0002694946, 0.0002694946 (x, y)
## extent      : -121.807, -121.725, 38.00413, 38.07204 (xmin, xmax, ymin, ymax)
## crs         : +proj=longlat +datum=WGS84 +no_defs +ellps=WGS84 +towgs84=0,0,0
## source      : memory
## names       : layer
## values      : 1, 10 (min, max)
```

We can see that knr is a RasterLayer but we do not know which cluster (1-10) belongs what LULC class. You can find that out by plotting them side-by-side with other reference layer and using unique color for each cluster.

```
# Define a color vector for 10 clusters (learn more about setting the color later)
mycolor <- c("#fef65b", "#ff0000", "#daa520", "#0000ff", "#0000ff", "#00ff00", "#cbb5b5",
             "#c3ff5b", "#ff7373", "#00ff00", "#808080")

par(mfrow = c(1,2))
plot(ndvi, col = rev(terrain.colors(10)), main = 'Landsat-NDVI')
plot(knr, main = 'Unsupervised classification', col = mycolor )
```



While for other purposes it is usually better to define more classes (and possibly fuse classes later), a simple classification like this one could be useful, e.g., merge cluster 4 and 5 to construct a water mask for the year 2011.

**Important** These 10 clusters are not useful unless you re-label them with LULC information class. You can save the raster output of clustering and open it in GIS platform (e.g. [QGIS](#)) and assign classes to each clusters. Multiple clusters can have the same class! A tutorial can be found [here](#). You can skip step 1 in the tutorial which provides guidelines for generating clusters.

You can change the colors in my mycolor. Learn more about selecting colors in R [here](#) and [here](#).

**Exercise** Plot 3-band RGB of landsat5 for the subset (extent e) and result of kmeans clustering side-by-side and make a table of land-use land-cover labels for the clusters. E.g. cluster 4 and 5 are water.



## SUPERVISED IMAGE CLASSIFICATION

Here we explore supervised classification. Various supervised classification algorithms exist, and the choice of algorithm can affect the results. Here we explore two related algorithms (CART and RandomForest).

In supervised classification, we have prior knowledge about some of the land-cover types through, for example, field-work, reference spatial data or interpretation of high resolution imagery (such as available on Google maps). Specific sites in the study area that represent homogeneous examples of these known land-cover types are identified. These areas are commonly referred to as training sites because the spectral properties of these sites are used to train the classification algorithm.

The following examples uses a Classification and Regression Trees (CART) classifier (Breiman et al. 1984) ([further reading](#)) to predict land use land cover classes in the study area.

We will perform the following steps:

- Generate sample sites based on a reference raster
- Extract cell values from Landsat data for the sample sites
- Train the classifier using training samples
- Classify the Landsat data using the trained model
- Evaluate the accuracy of the model

### 5.1 Reference data

The [National Land Cover Database 2011 \(NLCD 2011\)](#) is a land cover product for the USA. NLCD is a 30-m Landsat-based land cover database spanning 4 epochs (1992, 2001, 2006 and 2011). NLCD 2011 is based primarily on a decision-tree classification of circa 2011 Landsat data.

You can find the class names in NLCD 2011 ([here](https://www.mrlc.gov/nlcd11_leg.php)) [[https://www.mrlc.gov/nlcd11\\_leg.php](https://www.mrlc.gov/nlcd11_leg.php)]. It has two pairs of class-values and classnames that correspond to the levels of land use and land cover classification system. These levels usually represent the level of complexity, level I being the simplest with broad land use land cover categories. Read [this report by Anderson et al](#) to learn more about this land use and land cover classification system for use with Remote Sensor data.

```
library(raster)
nlcd <- brick('data/rs/nlcd-L1.tif')
names(nlcd) <- c("nlcd2001", "nlcd2011")

# The class names and colors for plotting
nlcdclass <- c("Water", "Developed", "Barren", "Forest", "Shrubland", "Herbaceous",
  ↪ "Planted/Cultivated", "Wetlands")
classdf <- data.frame(classvalue1 = c(1,2,3,4,5,7,8,9), classnames1 = nlcdclass)
```

(continues on next page)

(continued from previous page)

```
# Hex codes of colors
classcolor <- c("#5475A8", "#B50000", "#D2CDC0", "#38814E", "#AF963C", "#D1D182", "
↪ #F8766D", "#C8E6F8")

# Now we ratify (RAT = "Raster Attribute Table") the nlcd2011 (define RasterLayer as ↪
↪ a categorical variable). This is helpful for plotting.
nlcd2011 <- nlcd[[2]]
nlcd2011 <- ratify(nlcd2011)
rat <- levels(nlcd2011)[[1]]

#
rat$landcover <- nlcdclass
levels(nlcd2011) <- rat
```

We did a lot of things here. Take a step back and read more about `ratify`.

**Note** There is no class with value 6.

## 5.2 Generate sample sites

As we discussed in the class, training and/or validation data can come from a variety of sources. In this example we will generate the training and validation sample sites using the NLCD reference RasterLayer. Alternatively, you can use predefined sites that you may have collected from other sources. We will generate the sample sites following a stratified random sampling to ensure samples from each LULC class.

```
# Load the training sites locations
# Set the random number generator to reproduce the results
set.seed(99)

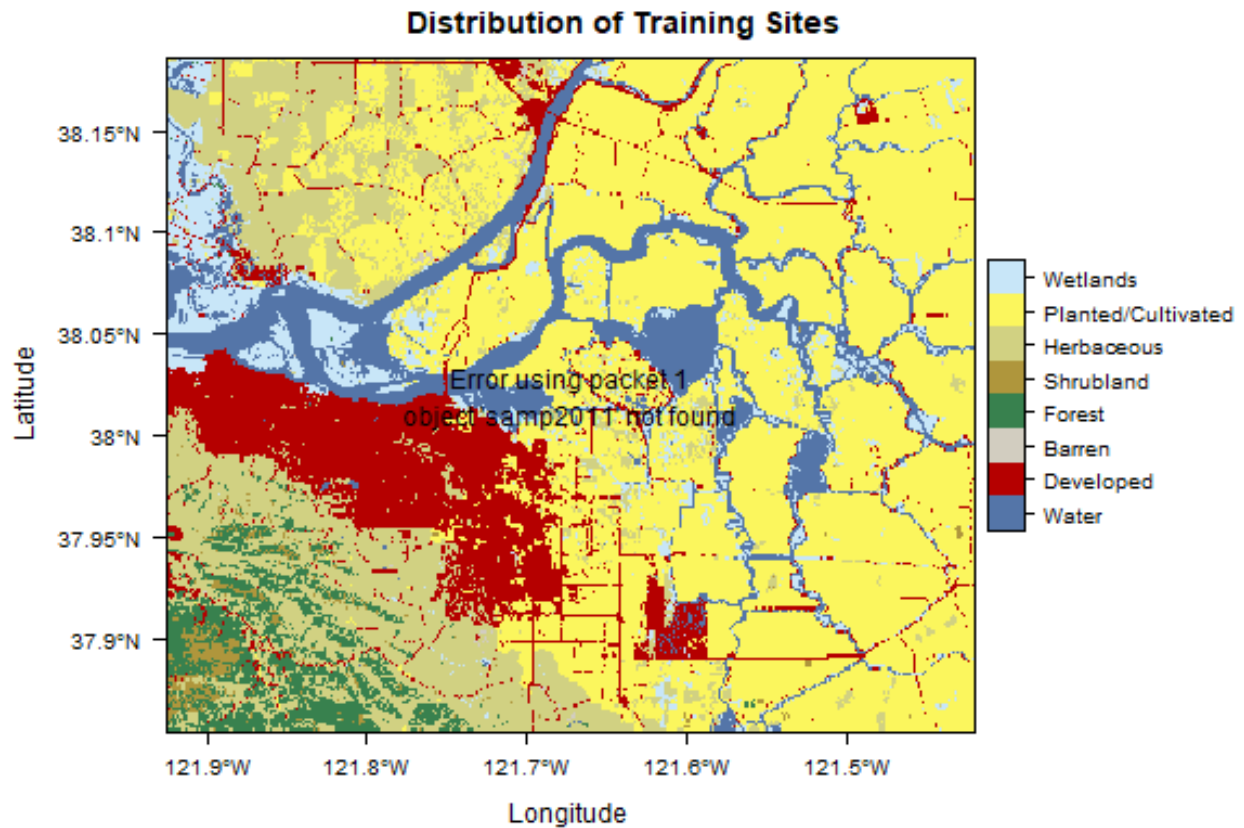
# Sampling
samp2011 <- sampleStratified(nlcd2011, size = 200, na.rm = TRUE, sp = TRUE)
samp2011
## class      : SpatialPointsDataFrame
## features    : 1600
## extent      : -121.9257, -121.4207, 37.85415, 38.18536 (xmin, xmax, ymin, ymax)
## crs         : +proj=longlat +datum=WGS84 +no_defs +ellps=WGS84 +towgs84=0,0,0
## variables   : 2
## names       : cell, nlcd2011
## min values  :      829,      1
## max values  : 2308569,      9
# Number of samples in each class
table(samp2011$nlcd2011)
##
##  1  2  3  4  5  7  8  9
## 200 200 200 200 200 200 200 200
```

You can see there are two variables in `samp2011`. The `cell` column contains cell numbers of `nlcd2011` sampled. `nlcd2011` column contains the class values (1-9). We will drop the `cell` column later.

Here `nlcd` has integer values between 1-9. You will often find classnames are provided as string labels (e.g. water, crop, vegetation). You will need to 'relabel' class names to integer or factors if only string labels are supplied before using them as response variable in the classification. There are several approaches that could be used to convert these classes to integer codes. We can make a function that will reclassify the character strings representing land cover classes into integers based on the existing factor levels.

Let's plot the training sites over the `nlcd2011` RasterLayer to visualize the distribution of sampling locations.

```
library(rasterVis)
plt <- levelplot(nlcd2011, col.regions = classcolor, main = 'Distribution of Training
↪Sites')
print(plt + layer(sp.points(samp2011, pch = 3, cex = 0.5, col = 1)))
```



rasterVis offers more advanced (trellis/lattice) plotting of Raster\* objects. Please install the package if it is not available for your machine.

### 5.3 Extract cell values for the sample sites

```
landsat5 <- stack('data/rs/centralvalley-2011LT5.tif')
names(landsat5) <- c('blue', 'green', 'red', 'NIR', 'SWIR1', 'SWIR2')
```

Once we have the sites, we can extract the cell values from `landsat5` RasterStack. These band values will be the predictor variables and “classvalues” from `nlcd2011` will be the response variable.

```
# Extract the layer values for the locations
sampvals <- extract(landsat5, samp2011, df = TRUE)

# sampvals no longer has the spatial information. To keep the spatial information you
↪ use `sp = TRUE` argument in `extract` function.

# drop the ID column
sampvals <- sampvals[, -1]

# combine the class information with extracted values
sampdata <- data.frame(classvalue = samp2011@data$nlcd2011, sampvals)
```

## 5.4 Train the classifier using training samples

Now we will train the classification algorithm using `training2011` dataset.

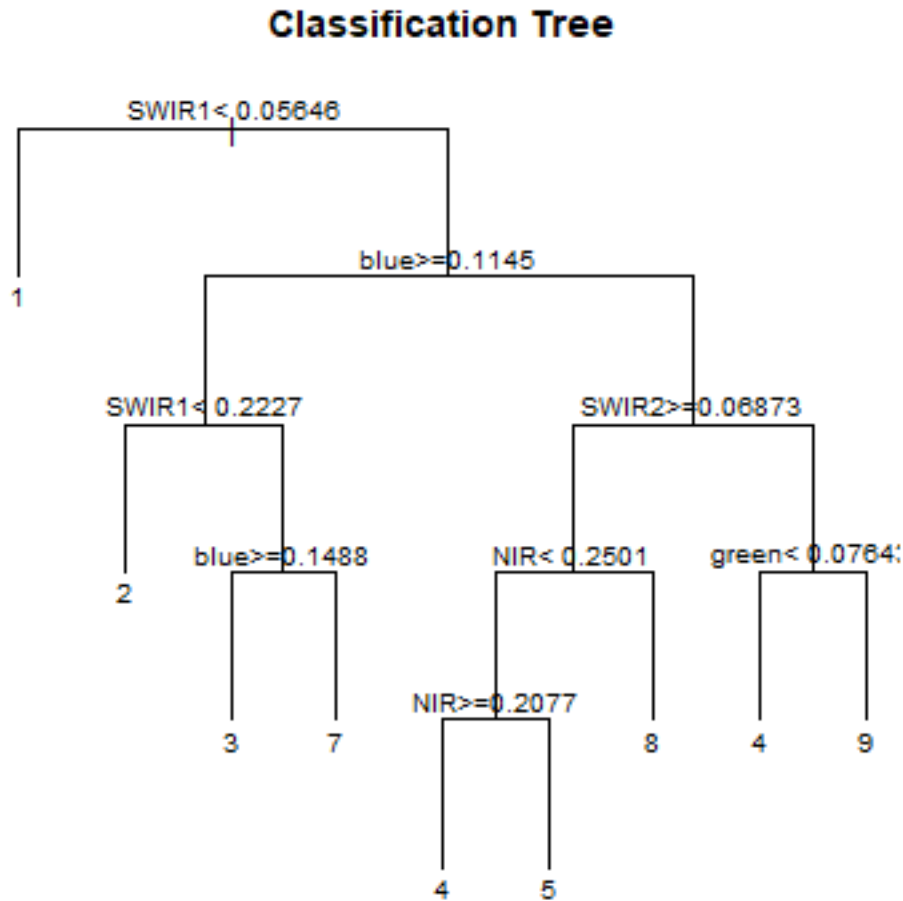
```
library(rpart)

# Train the model
cart <- rpart(as.factor(classvalue)~., data=sampdata, method = 'class', minsplit = 5)

# print(model.class)

# Much cleaner way is to plot the trained classification tree
plot(cart, uniform=TRUE, main="Classification Tree")
text(cart, cex = 0.8)
```





In the classification tree plot classvalues are printed at the leaf node. You can find the corresponding land use land cover names from the `classdf` dataframe.

See `?rpart.control` to set different parameters for building the model.

You can print/plot more about the `cart` model created in the previous example. E.g. you can use `plotcp(cart)` to learn about the cost-complexity (`cp` argument in `rpart`).

## 5.5 Classify

Now we have our trained classification model (`cart`), we can use it to make predictions, that is, to classify all cells in the `landsat5` RasterStack.

**Important** The names in the Raster object to be classified should exactly match those expected by the model. This will be the case if the same Raster object was used (via `extract`) to obtain the values to fit the model (previous example).

```
# Now predict the subset data based on the model; prediction for entire area takes_
↪ longer time
pr2011 <- predict(landsat5, cart, type='class')
pr2011
## class      : RasterLayer
## dimensions : 1230, 1877, 2308710  (nrow, ncol, ncell)
## resolution : 0.0002694946, 0.0002694946  (x, y)
## extent     : -121.9258, -121.42, 37.85402, 38.1855  (xmin, xmax, ymin, ymax)
## crs        : +proj=longlat +datum=WGS84 +no_defs +ellps=WGS84 +towgs84=0,0,0
## source     : memory
## names      : layer
## values     : 1, 9  (min, max)
## attributes :
##           ID value
## from: 1      1
## to  : 8      9
```

Now plot the classification result using `rasterVis`. See will set the classnames for the classvalues.

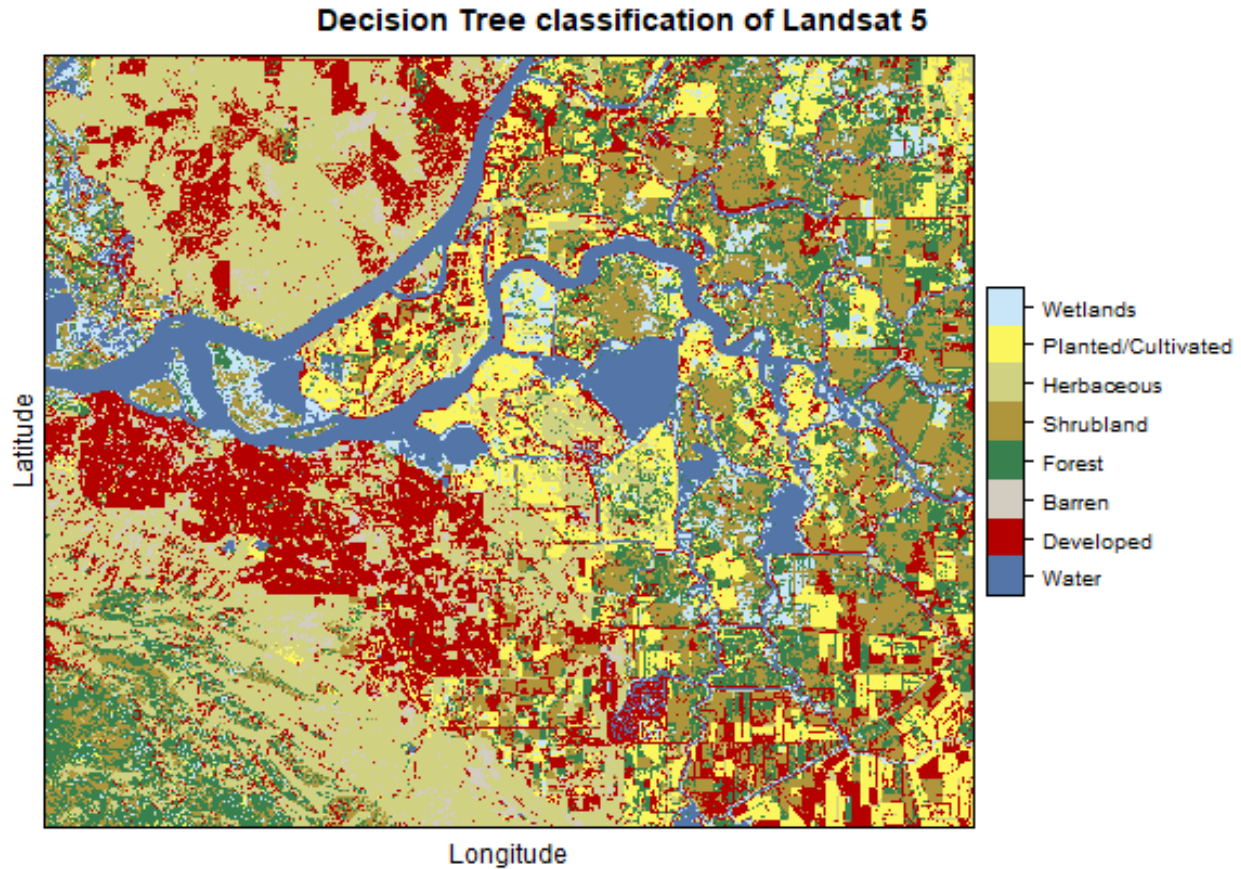
```
pr2011 <- ratify(pr2011)

rat <- levels(pr2011)[[1]]

rat$legend <- classdf$classnames

levels(pr2011) <- rat

levelplot(pr2011, maxpixels = 1e6,
          col.regions = classcolor,
          scales=list(draw=FALSE),
          main = "Decision Tree classification of Landsat 5")
```



**Exercise 1** Plot `nlcd2011` and `pr2011` side-by-side and comment about the accuracy of the prediction (e.g. mixing between cultivated crops, pasture, grassland and shrubs).

You may need to select more samples and use additional predictor variables. The choice of classifier also plays an important role.

## 5.6 Model evaluation

Now let's assess the accuracy of the model to get an idea of how accurate the classified map might be. Two widely used measures in remote sensing community are "overall accuracy" and "kappa". You can perform the accuracy assessment using the independent samples (`validation2011`).

To evaluate any model, you can use k-fold cross-validation. In this technique the data used to fit the model is split into  $k$  groups (typically 5 groups). In turn, one of the groups will be used for model testing, while the rest of the data is

used for model training (fitting).

```
library(dismo)
set.seed(99)
j <- kfold(sampdata, k = 5, by=sampdata$classvalue)
table(j)
## j
##   1    2    3    4    5
## 320 320 320 320 320
```

Now we train and test the model five times, each time computing a confusion matrix that we store in a list.

```
x <- list()

for (k in 1:5) {
  train <- sampdata[j!= k, ]
  test  <- sampdata[j == k, ]
  cart  <- rpart(as.factor(classvalue)~., data=train, method = 'class', minsplit = 5)
  pclass <- predict(cart, test, type='class')
  # create a data.frame using the reference and prediction
  x[[k]] <- cbind(test$classvalue, as.integer(pclass))
}
```

Now combine the five list elements into a single data.frame, using `do.call` and compute a confusion matrix.

```
y <- do.call(rbind, x)
y <- data.frame(y)
colnames(y) <- c('observed', 'predicted')
```

```
conmat <- table(y)
# change the name of the classes
colnames(conmat) <- classdf$classnames
rownames(conmat) <- classdf$classnames
conmat
```

	predicted					
observed	Water	Developed	Barren	Forest	Shrubland	Herbaceous
Water	172	6	1	3	3	1
Developed	8	93	22	10	8	37
Barren	6	54	52	6	8	62
Forest	0	3	2	102	49	8
Shrubland	0	8	1	52	105	12
Herbaceous	1	16	24	4	16	121
Planted/Cultivated	2	20	5	24	37	24
Wetlands	14	12	1	32	25	16

```
##
## predicted
## observed      Planted/Cultivated Wetlands
## Water                1      13
## Developed           20      2
## Barren              5      7
## Forest             13     23
## Shrubland          17      5
## Herbaceous         17      1
## Planted/Cultivated  81      7
## Wetlands           32     68
```

**Exercise 2** Comment on the miss-classification between different classes.

**Exercise 3** Can you think of ways to improve the accuracy.

Compute overall accuracy and Kappa statistics.

Overall accuracy

```
# number of cases
n <- sum(conmat)
n
## [1] 1600

# number of correctly classified cases per class
diag <- diag(conmat)

# Overall Accuracy
OA <- sum(diag) / n
OA
## [1] 0.49625
```

Kappa

```
# observed (true) cases per class
rowsums <- apply(conmat, 1, sum)
p <- rowsums / n

# predicted cases per class
colsums <- apply(conmat, 2, sum)
q <- colsums / n

expAccuracy <- sum(p*q)
kappa <- (OA - expAccuracy) / (1 - expAccuracy)
kappa
## [1] 0.4242857
```

Producer and user accuracy

```
# Producer accuracy
PA <- diag / colsums

# User accuracy
UA <- diag / rowsums

outAcc <- data.frame(producerAccuracy = PA, userAccuracy = UA)

outAcc
##               producerAccuracy userAccuracy
## Water                0.8472906         0.860
## Developed            0.4386792         0.465
## Barren                0.4814815         0.260
## Forest               0.4377682         0.510
## Shrubland            0.4183267         0.525
## Herbaceous           0.4306050         0.605
## Planted/Cultivated   0.4354839         0.405
## Wetlands             0.5396825         0.340
```

**Exercise 4** Perform the classification using Random Forest classifiers from the package `randomForest` <<https://cran.r-project.org/web/packages/randomForest/index.html>>‘\_\_’. For help see this [discussion](#).

**Exercise 5** Plot the results of `rpart` and `rf` classifier side-by-side.

**Exercise 6 (optional)** Please repeat the steps for the year 2001 using `rpart`. You will use cloud-free composite image from Landsat 7 with 6-bands. The reference data will be the [National Land Cover Database 2001 \(NLCD 2001\)](#)

for the subset of the Central Valley regions.

**Exercise 7 (optional)** You have been training your classifiers using 160 samples for each class. Now we want to see the effect of sample size on classification. Please repeat the steps with different subset of the sample size (e.g. 120, 100). Use the same holdout samples for accuracy assessment.